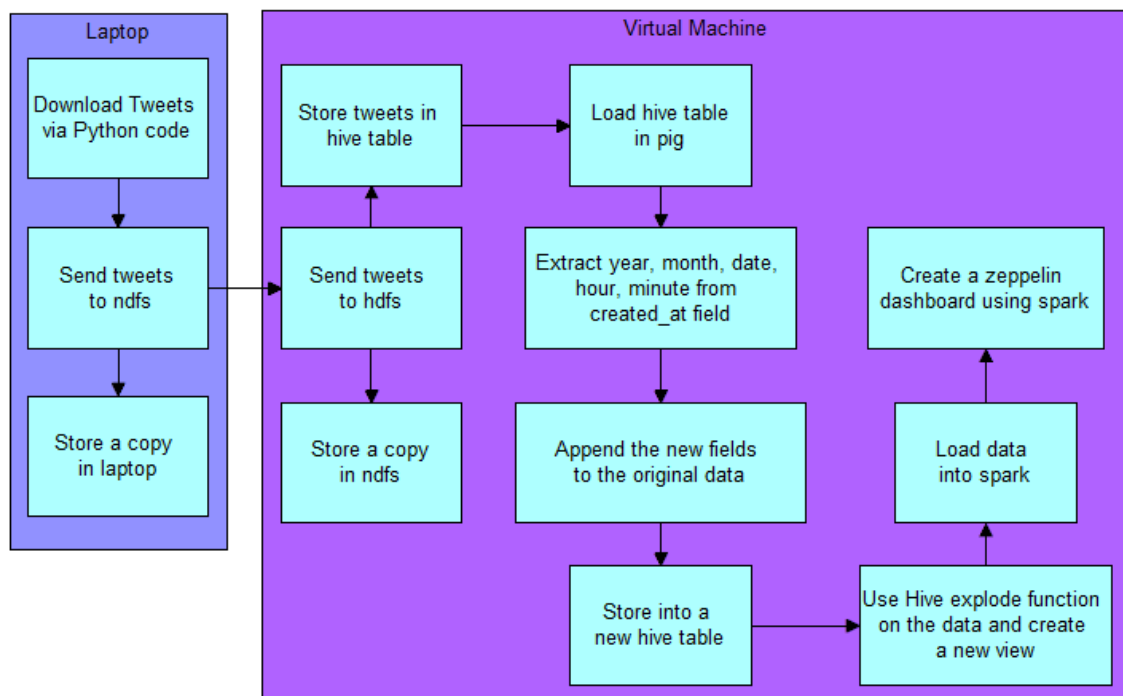


1. Project Flow

The project in this course is to make an Extract, Transform, Load (ETL) pipeline which touches all the tools taught in the course so that you can have a full hands-on experience on all of them. Keep in mind that the ETL pipeline proposed for the project is a highly inefficient one and is only given so that you can have some practice.

This document contains all the commands for the project along with their usage to help you in your project.

1.1. ETL pipeline



A high level overview of the ETL pipeline is given in the above diagram. We will use a python program to fetch tweets from Twitter using the tweepy API every hour. A new file will be created by the python program which will be stored on the laptop and a copy will be sent to the ndfs of the sandbox. A copy of the file will then be sent to the HDFS and stored there as well.

Then we will load the file in the Hive table created using Hive's native JSON SERDE (**external SERDE can be used as well**) which will parse the tweets and extract the data required. Next, the hive data will be loaded into Pig where we will break the CREATED_AT column into year, month, date, hour and minute columns, append them to the text and id columns and store the result in a new hive table.

After this we will use the remove any new line character from the text column and use the Hive explode function on the text column. Furthermore, we will load the data into Spark using the zeppelin notebook and create a dashboard from the data.

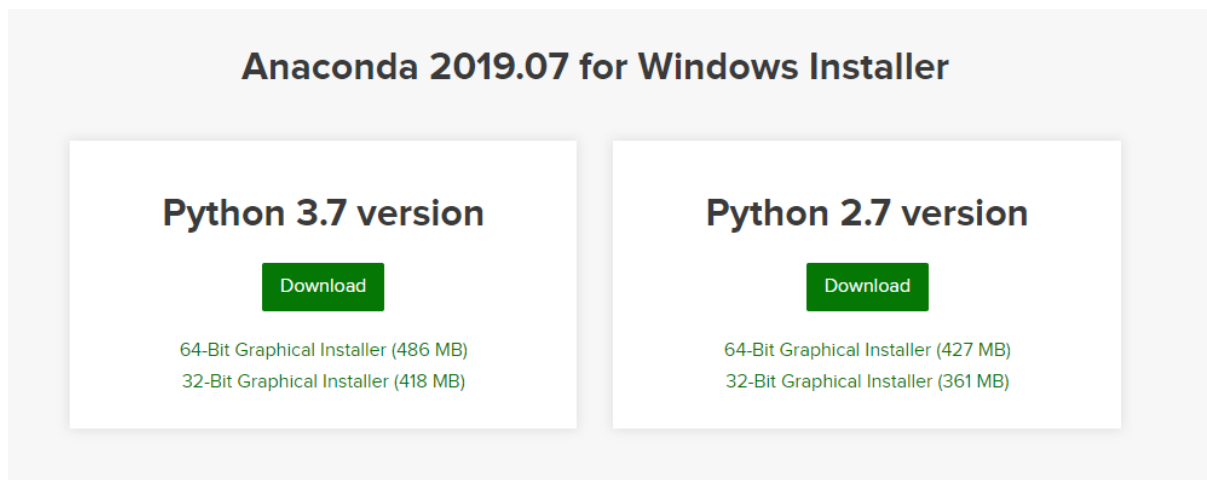
2. Pre-requisites

2.1. Installing Anaconda

Anaconda is a complete package for Python which provides a stable version of Python and a complete IDE and notebook for python coding. In addition, it also provides tools for debugging and testing making it a go-to application for python development. You will be provided with the Anaconda 5.0.1 executable for Windows 64-bit systems or you can download it directly from downloads section on Anaconda's website. **The python code provided in a separate .py file.**

Here's the download link for Anaconda's website. Choose Python 3.7 installer according to your machine specs (32/64 bit).

<https://www.anaconda.com/download/>

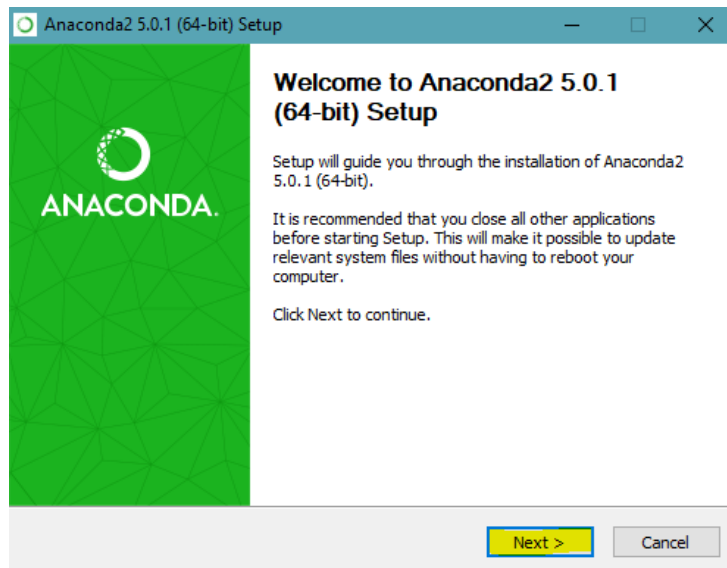


Anaconda 2019.07 for Windows Installer

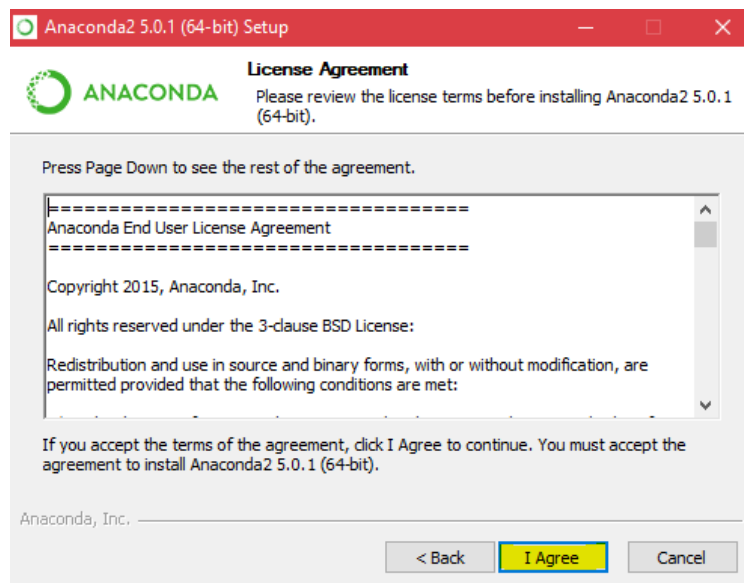
Python 3.7 version	Python 2.7 version
Download	Download
64-Bit Graphical Installer (486 MB)	64-Bit Graphical Installer (427 MB)
32-Bit Graphical Installer (418 MB)	32-Bit Graphical Installer (361 MB)

Follow the installation wizard and install Anaconda. Make sure to follow the process described below or there can be problems with your installation:

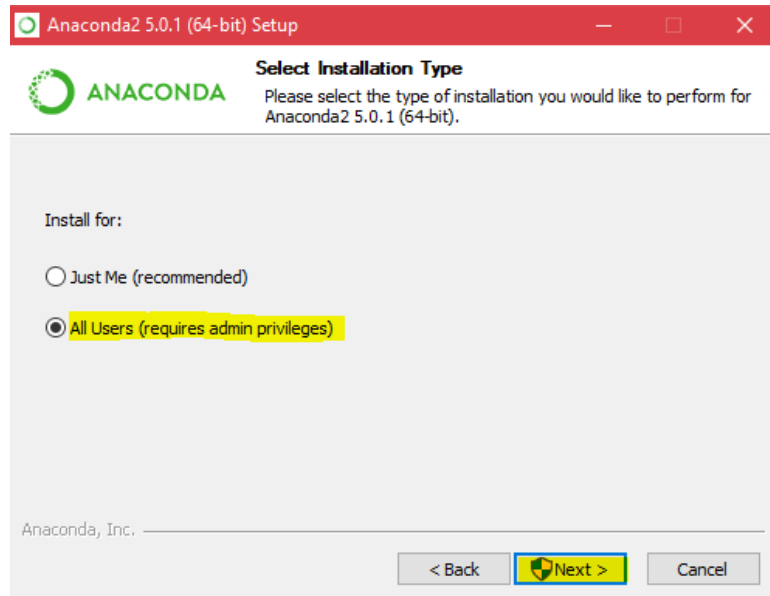
- Open the installer and click Next



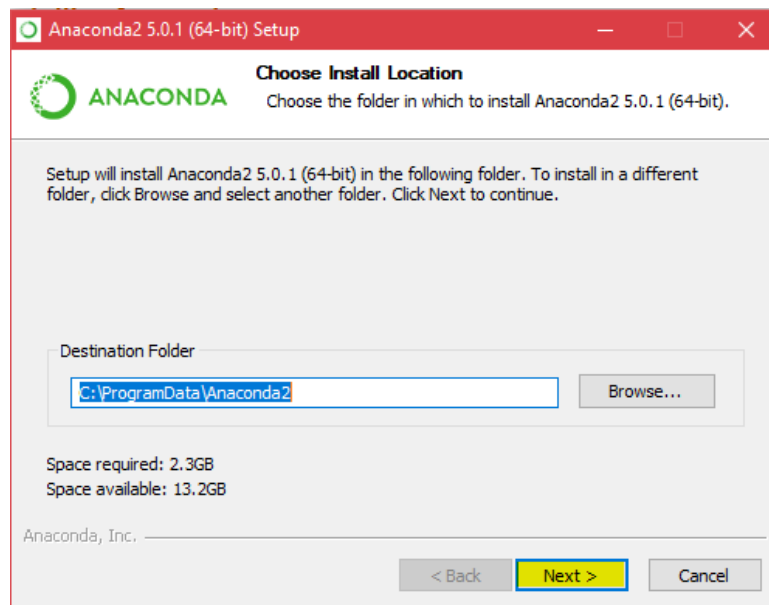
- Click on I Agree



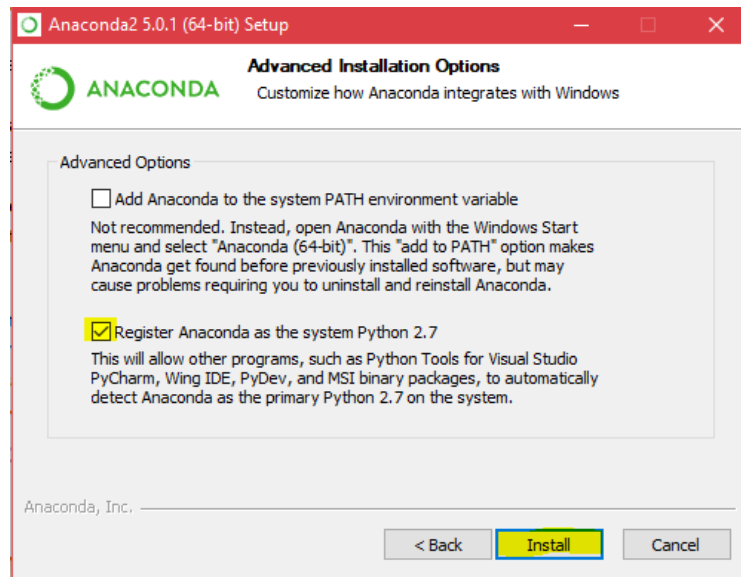
- Select All Users installation type



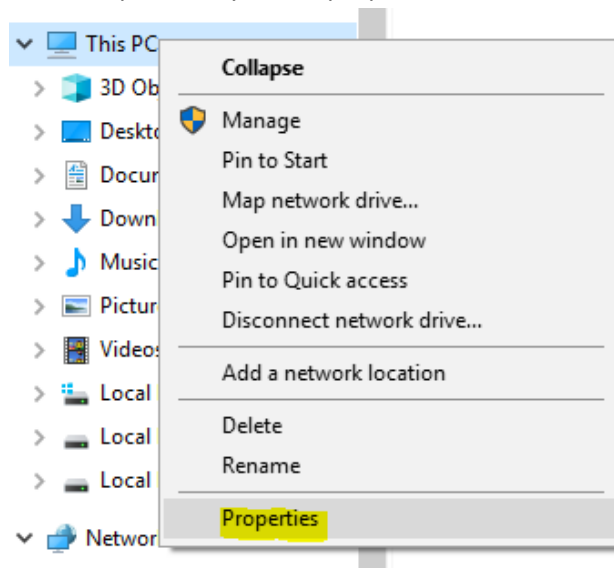
- Select the folder you want to install Anaconda in and click next. **Take note of the installation path entered in the Destination Folder as we will need it later on.** Its best to copy the path in a notepad file for later use.



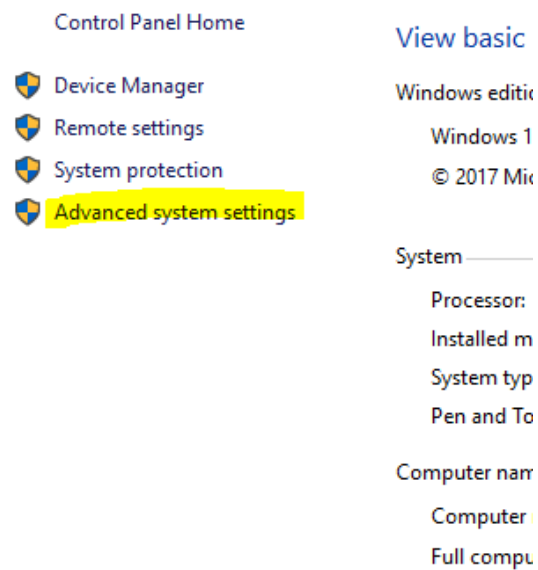
- In the Advanced options step for the make sure the Add Anaconda to my PATH environment variable is unchecked as below. This is an option but as the warning below it states, it can cause problems. Check the other option and click install



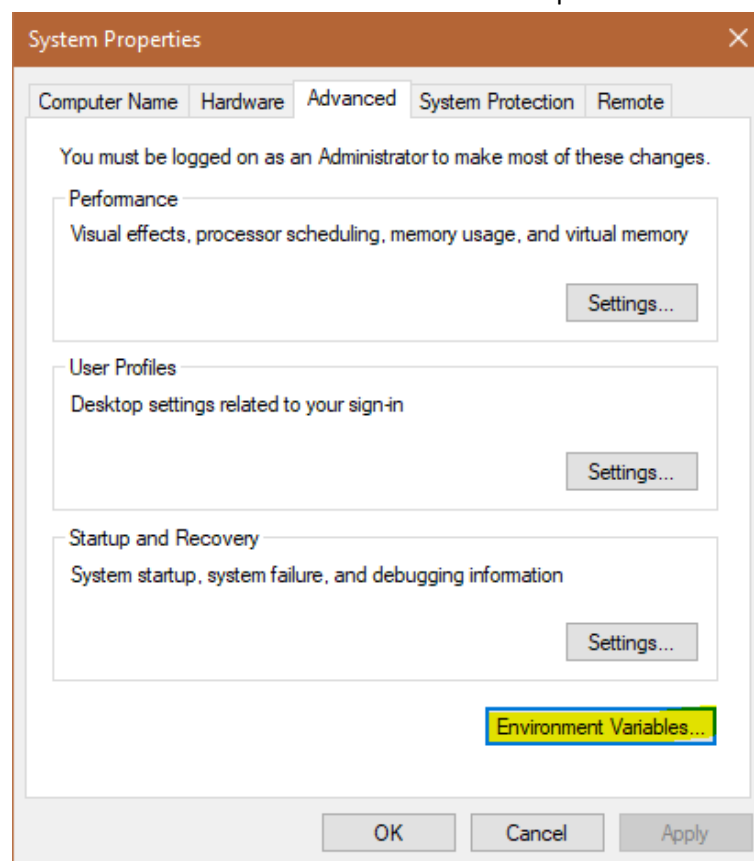
- The installation will be done and just click finish. Now we need to add the path for Anaconda ourselves as we unchecked the option in the previous point. Open file explorer and on the left side right click the computer name (usually named This PC but can be different). This will open the system's properties.



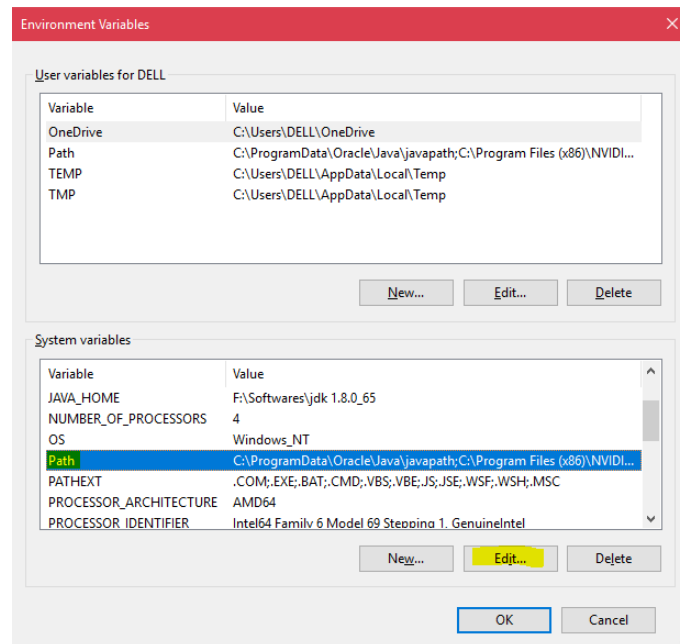
- Click on Advanced System Settings on the left hand side of the system properties window. A new window will open.



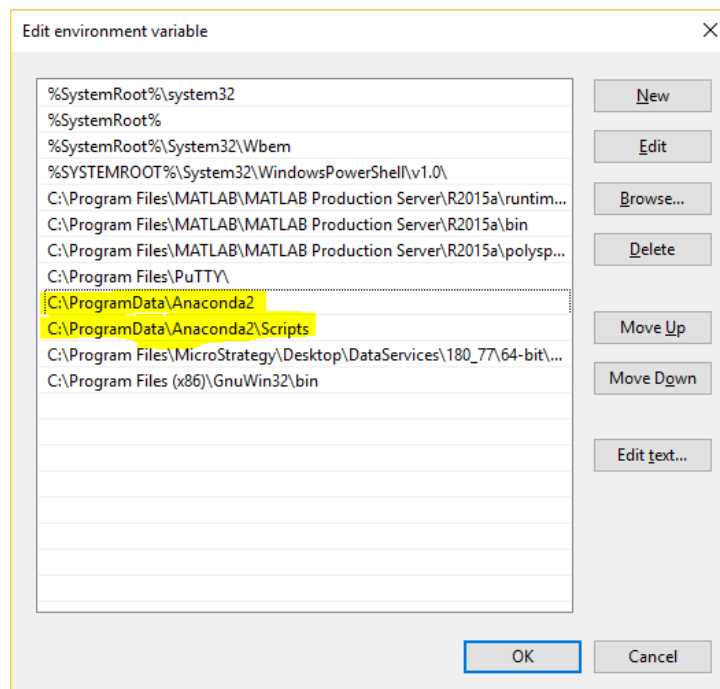
- Click on Environment Variables button. A new window will open.



- Find the Path variable in the System variables tab, select it and click Edit. A new window will open.



- Click New button and paste the path for Anaconda we saved earlier which is usually (**C:\ProgramData\Anaconda2**). Make sure to add (**C:\ProgramData\Anaconda2\scripts**) into variables as well (as shown below).
- **Make sure that you DON'T delete previously set environment variables.**
- Then click OK. Click OK again on the Environment Variables window and OK once again on the System Properties window. Now the path is set for Anaconda.



2.2. Installing tweepy

Next we need to install the tweepy library for Python. Tweepy provides a twitter API for Python and we can use it to fetch public tweets from Twitter.

- Open the Anaconda Prompt. Enter the command
where python
to make sure python is installed correctly. The command should return the path of the python installation. Next enter the command
pip install tweepy
to install tweepy. Pip will automatically find and install tweepy.

```
(C:\ProgramData\Anaconda2) C:\Users\DELL>where python
C:\ProgramData\Anaconda2\python.exe

(C:\ProgramData\Anaconda2) C:\Users\DELL>pip install tweepy
Collecting tweepy
  Downloading https://files.pythonhosted.org/packages/05/f1/2e8c7b202dd04117a378
ac0c55cc7dafa80280ebd7f692f1fa8f27fd6288/tweepy-3.6.0-py2.py3-none-any.whl
Collecting requests-oauthlib>=0.7.0 (from tweepy)
  Downloading https://files.pythonhosted.org/packages/77/34/d0957563f20b259a31c1
2f14e858d79f2e66eb539d3c1b9ab7077ef030ca/requests_oauthlib-0.8.0-py2.py3-none-an
y.whl
Requirement already satisfied: six>=1.10.0 in c:\programdata\anaconda2\lib\site-
packages (from tweepy)
Requirement already satisfied: requests>=2.11.1 in c:\programdata\anaconda2\lib\
site-packages (from tweepy)
Requirement already satisfied: PySocks>=1.5.7 in c:\programdata\anaconda2\lib\si
te-packages (from tweepy)
Collecting oauthlib>=0.6.2 (from requests-oauthlib>=0.7.0->tweepy)
  Downloading https://files.pythonhosted.org/packages/e0/ac/c6a0c98788aa0d619151
90d089e9ebe680905a94261effe3936eb8fe356f/oauthlib-2.0.7-py2.py3-none-any.whl (12
4kB)
    100% |#####| 133kB 208kB/s
Requirement already satisfied: chardet<3.1.0,>=3.0.2 in c:\programdata\anaconda2
\lib\site-packages (from requests>=2.11.1->tweepy)
Requirement already satisfied: idna<2.7,>=2.5 in c:\programdata\anaconda2\lib\si
te-packages (from requests>=2.11.1->tweepy)
Requirement already satisfied: urllib3<1.23,>=1.21.1 in c:\programdata\anaconda2
\lib\site-packages (from requests>=2.11.1->tweepy)
Requirement already satisfied: certifi>=2017.4.17 in c:\programdata\anaconda2\li
b\site-packages (from requests>=2.11.1->tweepy)
Installing collected packages: oauthlib, requests-oauthlib, tweepy
Successfully installed oauthlib-2.0.7 requests-oauthlib-0.8.0 tweepy-3.6.0
You are using pip version 9.0.1, however version 10.0.1 is available.
You should consider upgrading via the 'python -m pip install --upgrade pip' comm
and.

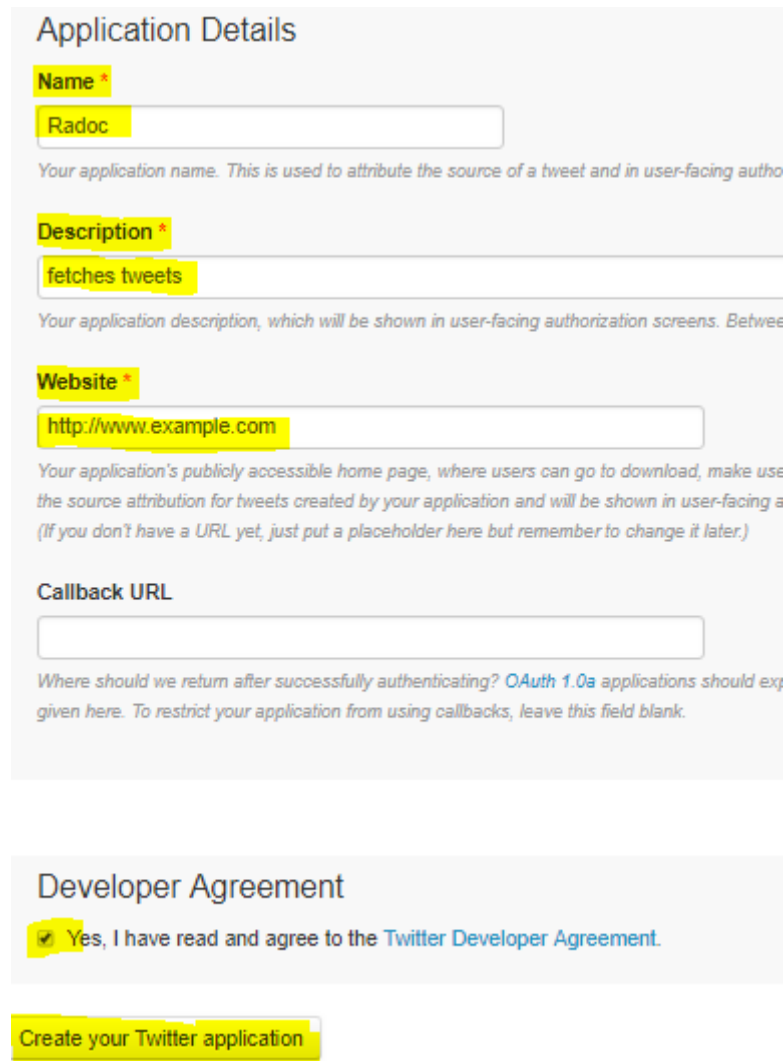
(C:\ProgramData\Anaconda2) C:\Users\DELL>
```


2.3. Creating a twitter app

If your account request rejected then you can skip this section and jump to section 3, as the API keys are also being provided within the python code.

Next we need to create a twitter app on the twitter website so that we can get the authorization credentials to access the site using tweepy.

- Go to <https://apps.twitter.com> and click on Create New App on the right-hand side. Enter the info required on the next page. Name has to be unique and the Website URL must be a valid URL but not necessarily an active site. It can be any website but the URL has to be valid. Agree to the Developer Agreement and click Create your Twitter application.



The screenshot shows the 'Application Details' form on the Twitter developer portal. The form includes fields for Name, Description, Website, and Callback URL, each with a placeholder text. Below the form is the 'Developer Agreement' section with a checkbox for agreement. At the bottom is a 'Create your Twitter application' button.

Application Details

Name *
Radoc
Your application name. This is used to attribute the source of a tweet and in user-facing authentication screens.

Description *
fetches tweets
Your application description, which will be shown in user-facing authorization screens. Between 1 and 140 characters.

Website *
<http://www.example.com>
Your application's publicly accessible home page, where users can go to download, make use of the source attribution for tweets created by your application and will be shown in user-facing authentication screens. (If you don't have a URL yet, just put a placeholder here but remember to change it later.)

Callback URL

Where should we return after successfully authenticating? OAuth 1.0a applications should expect to be redirected to this URL. To restrict your application from using callbacks, leave this field blank.

Developer Agreement
☒ Yes, I have read and agree to the [Twitter Developer Agreement](#).

Create your Twitter application

- Your app home will appear after the app has been created successfully. Click on the Keys and Access Tokens tab to view the keys required. On the top there will be the Consumer key

and the Consumer Secret. Copy these two keys for later use in the python file we will create later.

Radoc

[Test OAuth](#)[Details](#)[Settings](#)[Keys and Access Tokens](#)[Permissions](#)

Application Settings

Keep the "Consumer Secret" a secret. This key should never be human-readable in your application.

Consumer Key (API Key)	0jI80sGVXDwZm4Zm8KwRD4XIb
Consumer Secret (API Secret)	4PNYoWRLJ6XN5HPbpmPcj93hmRPWEwsfw4QIEv4QAhtEAXWnuX
Access Level	Read and write (modify app permissions)
Owner	faranshahid
Owner ID	231112714

- Scroll down and click on Generate Access Token to create your Access Token keys. Copy these two keys for later use as well.

Your Access Token

This access token can be used to make API requests on your own account's behalf. Do not share your access token secret with anyone.

Access Token	231112714-WROcnkqpdNhc0oRMgNb7enEcwgqkcFJrD4fc3Dcp
Access Token Secret	GhfGI1uYv04c4yUnV0x8XO6cHdZv8PwxI0MfPGFYICARX
Access Level	Read and write
Owner	faranshahid
Owner ID	231112714

Token Actions

[Regenerate My Access Token and Token Secret](#)[Revoke Token Access](#)

You can visit <https://apps.twitter.com> anytime to get these keys again. You can also generate new keys using the specified buttons.

3. ETL commands

3.1. Generating tweets

Note: Python code is provided as a separate .py file so please don't copy from here.

The data we will be feeding into the ETL pipeline will be tweets pertaining to a particular hashtag. These tweets are fetched through a python program which uses the tweepy library to fetch tweets. The complete code along with comments is given below.

The code below is for Python 2.7. To run this code on Python 3 only minor changes are required. Python is quite sensitive to tabs and space indentations in code so there will be errors when copying the code below. It is recommended to write the code yourself or use the python file if it is provided. The lines in light grey are comments to help you understand what the code does.

```
# coding: utf-8
# In[26]:
# importing the required libraries.
import tweepy
import json
import time
import datetime
# Authentication details. Enter your own twitter app keys here.
consumer_key = ""
consumer_secret = ""
access_key = ""
access_secret = ""
# Enter the hashtag you want to search for here.
accountvar = "Football"
# getting the current date and time.
t = datetime.datetime.now()
# sorting the acquired date and time into the format we want as Windows
# doesn't allow us to include : in file names.
a = t.strftime('%Y-%m-%d-%H-%M')
# specifying the output file name.
outputfilejson = accountvar+"_"+str(a)+".json"
# This is the listener, responsible for receiving data.
class StdOutListener(tweepy.StreamListener):
    # this class will be called before any of the others once the
    connection with the Twitter API is made.
    def __init__(self, time_limit=60):
        # setting current time as start time.
        self.start_time = time.time()
        # setting time limiter to pause stream. Current time limit is 60
seconds as specified in the function definition.
        self.limit = time_limit
        super(StdOutListener, self).__init__()
        # defining the on_data function. This will tell the compiler what to
do whenever new data is received.
    def on_data(self, data):
        # setting the time limit checker. It will let the stream fetch
```

```
data as long as the time limit has not been reached.
    if (time.time() - self.start_time) < self.limit:
        # loading the fetched encoded json tweet into the decoded
variable.
        decoded = json.loads(data)
        try:
            if isinstance(decoded, dict):
                # decoding the json tweet and loading it in decoded.
                decoded = json.dumps(decoded)
                # writing the decoded tweet into the output file.
                outfile.write(decoded)
                # adding a new line after the tweet into the output
file.

                outfile.write('\n')
                # printing the stuff we are writing into the output
file.

                print decoded+'\n'
            # handling exceptions in case there is some error
            except (NameError, KeyError, AttributeError):
                pass
            return True
        # code will go here once the time limit is reached. The following
lines will disconnect the stream.
        else:
            time.sleep(10)
            stream.disconnect()
            return False
        # code will go here if there is any error while making connection with
the Twitter API.
        def on_error(self, status):
            # printing the received error.
            print status
# specifying the main method. This method is always called first when the
code starts so we do all the initializing here.
if __name__ == '__main__':
    # calling the listener class.
    l = StdOutListener()
    # specifying the authorization handler of the Twitter API and giving
it the variables we defined at the start.
    auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
    # setting the access token with the tokens of our Twitter API.
    auth.set_access_token(access_key, access_secret)
    # creating a new output file and opening it in write back mode.
    with open(outputfilejson, 'wb') as outfile:
        print "Showing all new tweets for " + accountvar
```

```
# initializing the stream with the Twitter API authorization keys.
stream = tweepy.Stream(auth, l)
# searching for the required hashtag.
stream.filter(track=[str("#" + accountvar)])
```

3.2. Windows Scheduler

The tweet fetching and transfer has to be done every hour so we use a task scheduler to automate the process. It can be set up easily and made to run a .bat file.

NOTE: Read this entire section before following the steps

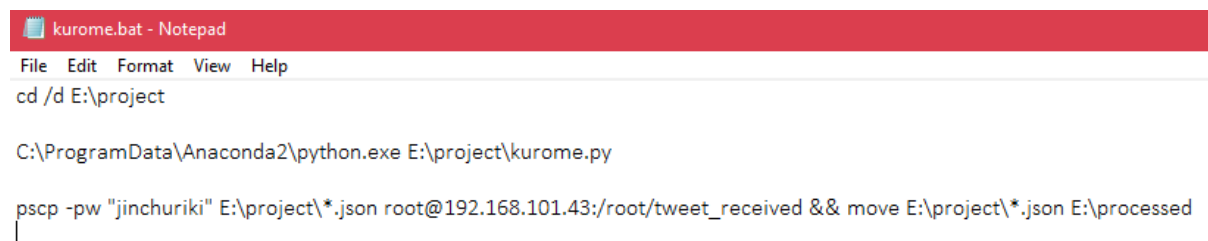
- Create a batch file in the folder where you placed your python file. Please make sure not to put them on the desktop. Open the batch file and paste the following code into it:

```
cd /d <path-of-your-python-file>
```

```
<path-to-your-python-executable>\python.exe <path-to-python-  
file>\<name-of-python-file.py>
```

```
pscp -pw "<your-vm-root-password>" <path-to-where-your-json-files-  
are-created>\*.json root@<your-vm-ip-address>:/root/<folder-name-  
where-tweets-are received> && move <path-to-where-your-json-files-  
are-created>\*.json <path-to-folder-where-json-files-are-to-be-  
stored-after-processing>
```

For example my code would be:






```
kurome.bat - Notepad
File Edit Format View Help
cd /d E:\project

C:\ProgramData\Anaconda2\python.exe E:\project\kurome.py

pscp -pw "jinchuriki" E:\project\*.json root@192.168.101.43:/root/tweet_received && move E:\project\*.json E:\processed
```


















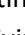



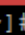


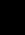
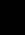
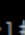

- Your json files will be created in the same folder as your python file or batch file. Create the folder where your json files will go after processing. The ETL pipeline on my laptop has been running for 3 months now so it looks like this:
E:\project is where I've placed my python and batch file and where my json files are created

This PC > Local Disk (E:) > project

Name	Date modified	Type	Size
 Football_2018-05-11-20-00.json	5/11/2018 8:01 PM	JSON File	92 KB
 kurome.py	5/3/2018 3:53 PM	Python File	4 KB
 kurome.bat	4/18/2018 8:23 PM	Windows Batch File	1 KB

E:\processed is where my json files are stored after processing

This PC > Local Disk (E:) > processed

Name	Date modified	Type	Size
 Football_2018-02-10-00-00.json	2/10/2018 12:01 AM	JSON File	122 KB
 Football_2018-02-10-01-00.json	2/10/2018 1:01 AM	JSON File	106 KB
 Football_2018-02-10-02-00.json	2/10/2018 2:01 AM	JSON File	69 KB
 Football_2018-02-10-03-00.json	2/10/2018 3:01 AM	JSON File	168 KB
 Football_2018-02-10-04-00.json	2/10/2018 4:01 AM	JSON File	177 KB
 Football_2018-02-10-05-00.json	2/10/2018 5:01 AM	JSON File	136 KB
 Football_2018-02-10-06-00.json	2/10/2018 6:01 AM	JSON File	193 KB
 Football_2018-02-10-07-00.json	2/10/2018 7:01 AM	JSON File	166 KB
 Football_2018-02-10-08-00.json	2/10/2018 8:01 AM	JSON File	126 KB
 Football_2018-02-10-09-00.json	2/10/2018 9:01 AM	JSON File	133 KB
 Football_2018-02-10-10-00.json	2/10/2018 10:01 AM	JSON File	134 KB
 Football_2018-02-10-11-00.json	2/10/2018 11:01 AM	JSON File	158 KB
 Football_2018-02-10-12-00.json	2/10/2018 12:01 PM	JSON File	187 KB
 Football_2018-02-10-14-00.json	2/10/2018 2:01 PM	JSON File	186 KB
 Football_2018-02-10-15-00.json	2/10/2018 3:01 PM	JSON File	179 KB
 Football_2018-02-10-16-00.json	2/10/2018 4:01 PM	JSON File	184 KB
 Football_2018-02-10-17-00.json	2/10/2018 5:01 PM	JSON File	210 KB
 Football_2018-02-10-18-00.json	2/10/2018 6:01 PM	JSON File	161 KB
 Football_2018-02-10-19-00.json	2/10/2018 7:01 PM	JSON File	226 KB
 Football_2018-02-10-20-00.json	2/10/2018 8:01 PM	JSON File	468 KB
 Football_2018-02-10-21-00.json	2/10/2018 9:01 PM	JSON File	245 KB
 Football_2018-02-10-22-00.json	2/10/2018 10:01 PM	JSON File	120 KB
 Football_2018-02-10-23-00.json	2/10/2018 11:01 PM	JSON File	232 KB
 Football_2018-02-11-00-00.json	2/11/2018 12:01 AM	JSON File	201 KB
 Football_2018-02-11-01-00.json	2/11/2018 1:01 AM	JSON File	162 KB
 Football_2018-02-11-02-00.json	2/11/2018 2:01 AM	JSON File	169 KB
 Football_2018-02-11-03-00.json	2/11/2018 3:01 AM	JSON File	162 KB
 Football_2018-02-11-04-00.json	2/11/2018 4:01 AM	JSON File	203 KB

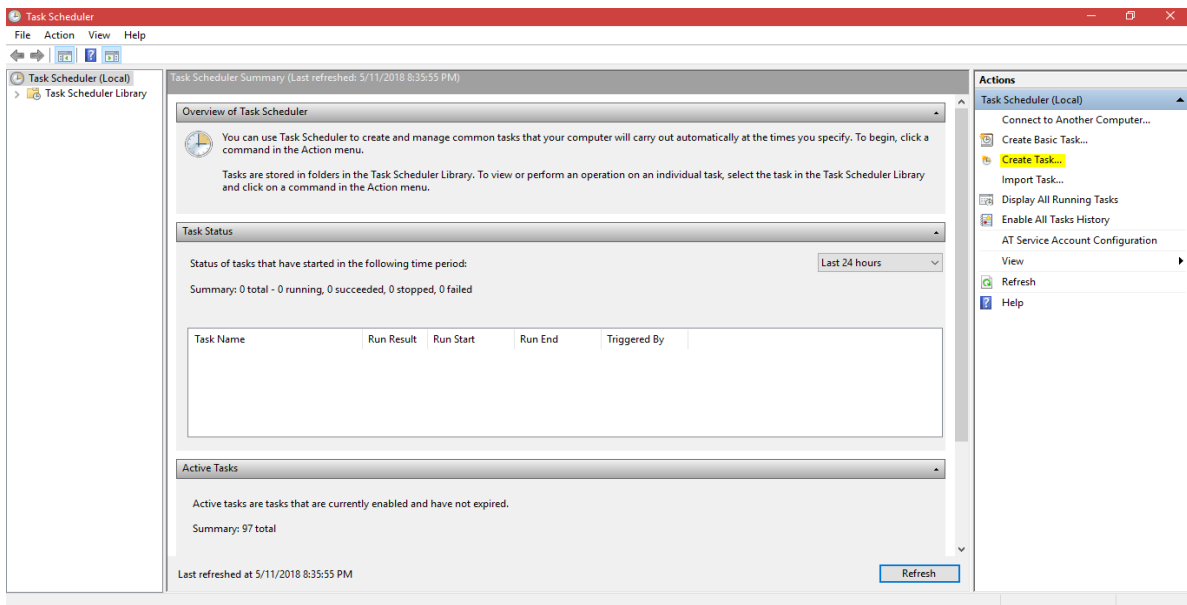
- After creating these two folders open the shell of your VM and create two folders there. One is for receiving json files from the laptop and the other one is for storing the json files after they have been sent from here to the hdfs (covered in the next section).

```

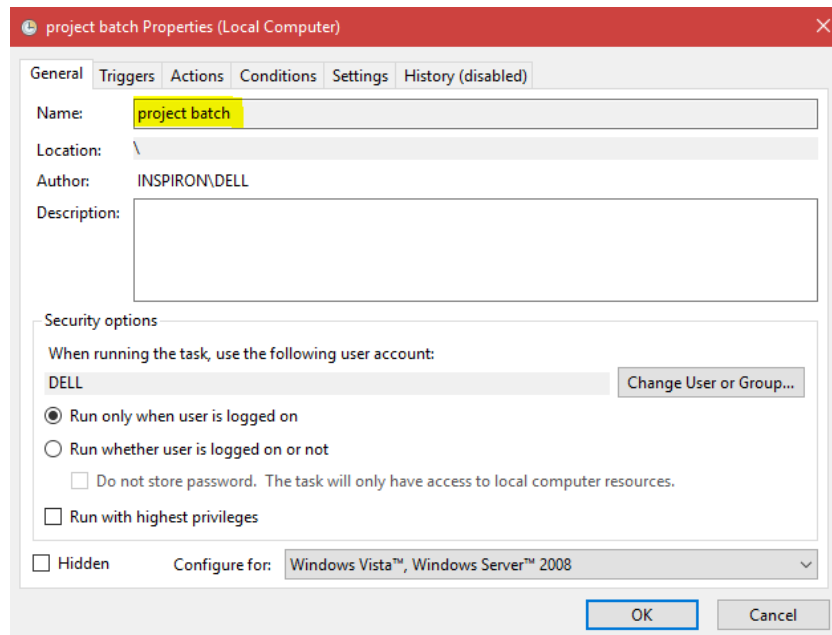
root@sandbox:~
[root@sandbox ~]# mkdir tweet_received
[root@sandbox ~]# mkdir tweet_processed
[root@sandbox ~]# ls
anaconda-ks.cfg  hiveql.hql          sandbox.info        tweet_processed
blueprint.json   id_rsa              start_ambari.sh    tweet_received
build.out        install.log         start_hbase.sh
derby.log        install.log.syslog  start_solr.sh
hi.json          proj_cron.sh        stop_solr.sh
[root@sandbox ~]#

```

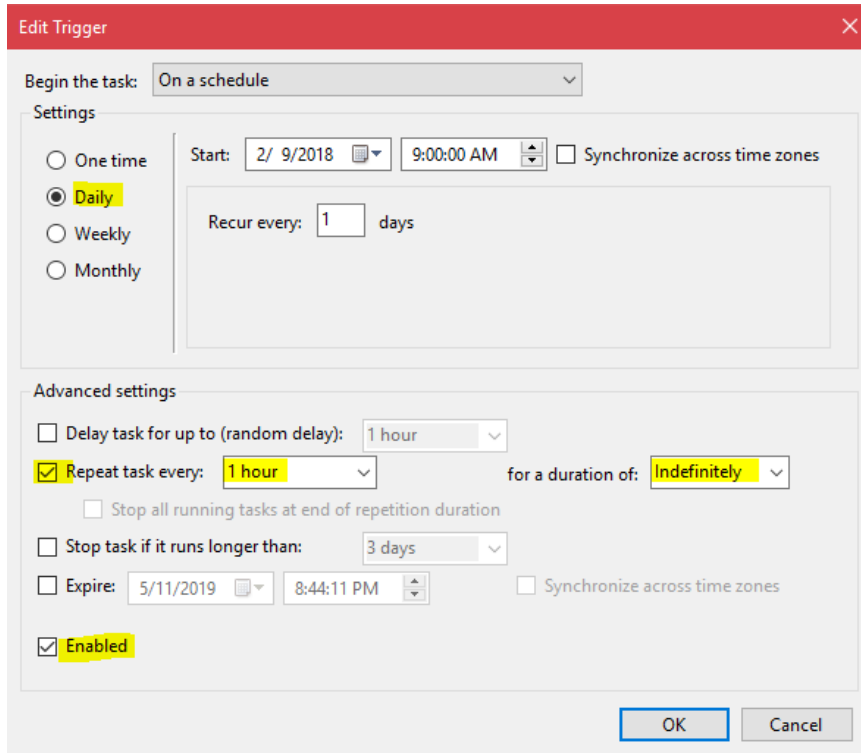
- Now get back to your laptop and open up Windows Task Scheduler. It is a default program of Windows and is already installed into your system. Click on Create Task on the right hand side of the window to create a new task.



- On the General tab enter the name you want to give to the task



- On the Triggers tab Click on New to create a new Trigger. Set scheduler to Daily, Repeat to one hour and indefinitely and check enabled as shown in the screenshot below. Click OK and the task will appear in the Triggers tab.



Edit Trigger

Begin the task: On a schedule

Settings

☐ One time
☒ **Daily**
☐ Weekly
☐ Monthly

Start: 2/ 9/2018 9:00:00 AM ☐ Synchronize across time zones

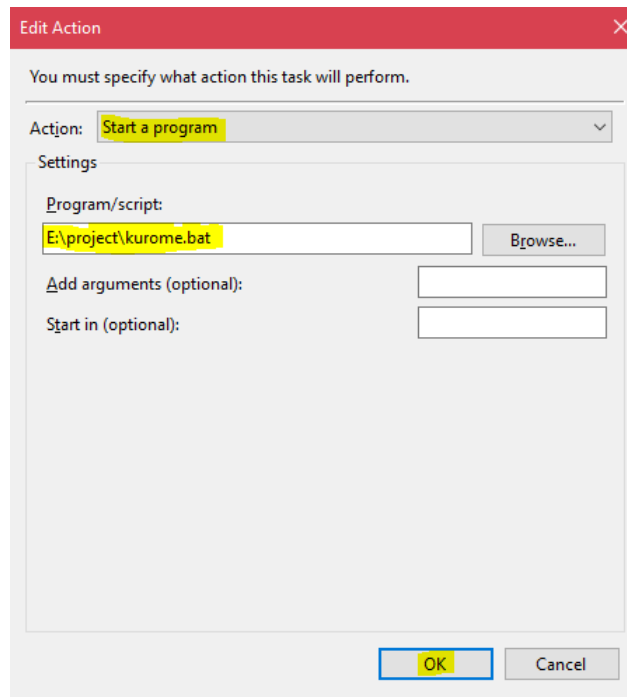
Recur every: 1 days

Advanced settings

☐ Delay task for up to (random delay): 1 hour
☒ **Repeat task every: 1 hour** for a duration of: **Indefinitely**
☐ Stop all running tasks at end of repetition duration
☐ Stop task if it runs longer than: 3 days
☐ Expire: 5/11/2019 8:44:11 PM ☐ Synchronize across time zones
☒ **Enabled**

OK Cancel

- Go to the Actions tab and Click New. Choose action as Start a Program and give the path to the batch file in the field specified. Click Ok to save it.



Edit Action

You must specify what action this task will perform.

Action: Start a program

Settings

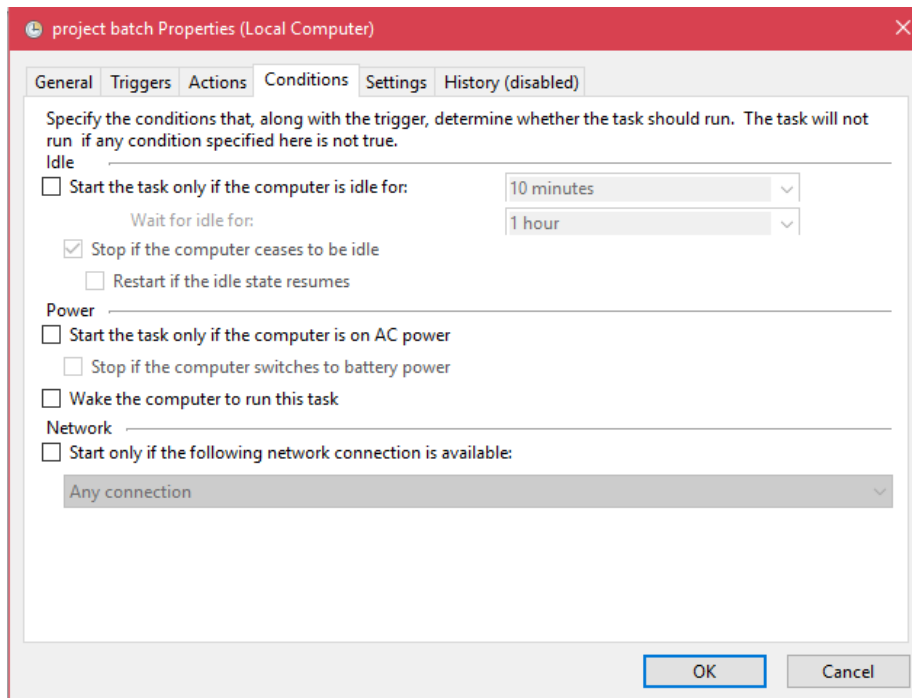
Program/script: E:\project\kurome.bat Browse...

Add arguments (optional):

Start in (optional):

OK Cancel

- On the Conditions tab uncheck everything. Click OK to create the Task



According to the time you have set in the Triggers tab the task will activate and run the batch file. Make sure the IP address in the batch file is updated every time the VM is started as the IP address changes every time.

3.3. CRON jobs scheduler

Now that our json files are moving from the laptop to NDFS we now need to move the file from ndfs to hdfs and for that we will use the built-in CentOS scheduler known as CRON jobs.

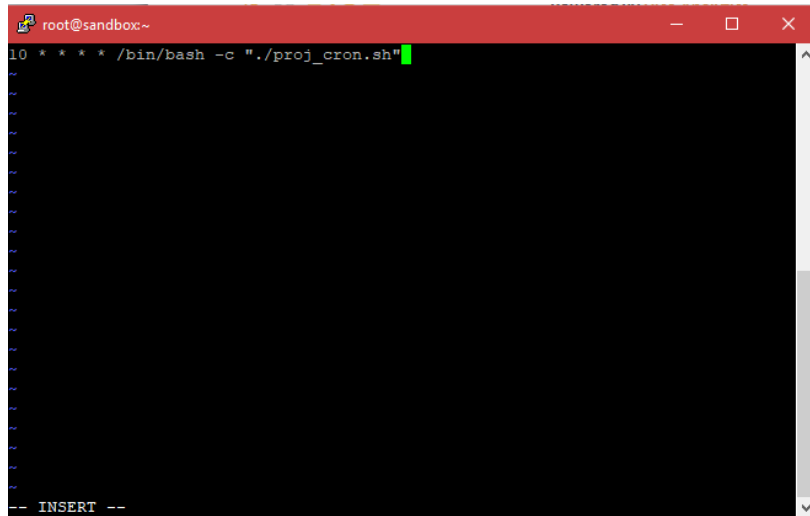
- On the shell enter the command `crontab -e` to edit the cron list and add a new cronjob.



```
root@sandbox:~  
[root@sandbox ~]# crontab -e
```

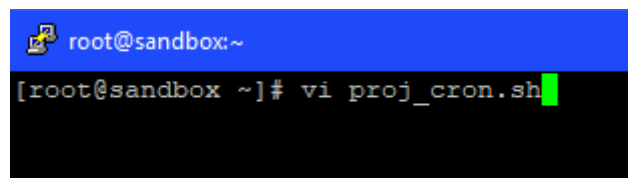
- The text editor will open up. Enter I to get into insert mode and add the following line:

```
10 * * * * /bin/bash -c "<your-desired-file-name>.sh"
```



```
root@sandbox:~  
10 * * * * /bin/bash -c "./proj_cron.sh"
```

- The above line shows that your scheduled job will run at 10th minute of every hour.
- Press `Esc` and enter `:wq` to save the cron job. To view the installed cron job you can enter `crontab -l` in the shell.
- Now we need to create the `.sh` file we specified in the cron job above. In the root directory enter
`vi <your-desired-file-name>.sh`



```
root@sandbox:~  
[root@sandbox ~]# vi proj_cron.sh
```

- A new text editor will open for editing the file. Press `I` to enter insert mode and add the following line:

```
hdfs dfs -put ./<folder-where-json-files-are-received-from-laptop>/*  
/<destination-folder-in-hdfs> && mv ~<folder-where-json-files-are-  
received-from-laptop>/* ~/<folder-to-put-files-after-processing>
```

[illegible]

- Press `Esc` and enter `:wq` to save the file. Now you need to give execution permissions to the file. Enter the following line in the shell for that:

```
chmod 777 <shell-file-name>.sh
```

```
root@sandbox:~  
[root@sandbox ~]# chmod 777 proj_cron.sh
```

- Now we need to create the folder in hdfs where the json files would be stored. We have already specified the folder name and location in the shell command above. Simply create a folder using Ambari GUI with the required name at the required location.

3.4. Add SERDE (Skip this section (3.4) if you're using Hive's native JSON SERDE)

At this point we are receiving the json files in HDFS and hence we have data to load. We now need to load this data into a hive table. The first step in loading the data is to get a serde for json. You are already provided with the serde needed in the Whatsapp group.

- Simply place the serde into a location in the HDFS using the Ambari GUI and give it full permissions.

usr

Name	Size	Last Modified	Owner	Group	Permission	Asc	Name
..							
employee.csv	0.3 kB	2018-05-03 20:46	admin	hdfs	-rw-r--r--		
hi.json	0.2 kB	2018-05-03 21:28	root	hdfs	-rw-r--r--		
id_rsa	1.6 kB	2018-05-03 19:44	root	hdfs	-rw-r--r--		
json-serde-1.1.9.9.jar	69.4 kB	2018-04-18 17:22	admin	hdfs	-rwxrwxrwx		
tweet_received	-	2018-05-12 01:10	admin	hdfs	-rwxr-xr-x		

Size

Group

1.3 kB

hdfs

1.2 kB

hdfs

1.6 kB

hdfs

69.4 kB

hdfs

hdfs

Read

Write

Execute

User

Read

Write

Execute

Group

Read

Write

Execute

Other

☐ Modify recursively

Close

Save changes

- Now go to Hive view and enter the following command to add the serde to the worksheet.

```
add jar hdfs:///<path-to-serde>/<name-of-serde>.jar;
```

Query Editor

Worksheet *

1 add jar hdfs:///usr/json-serde-1.1.9.9.jar;

NOTE: Remember that using this method the serde will only be added for the duration that the worksheet is active. Once the Hive view is closed the serde will automatically be removed and will have to be added next time.

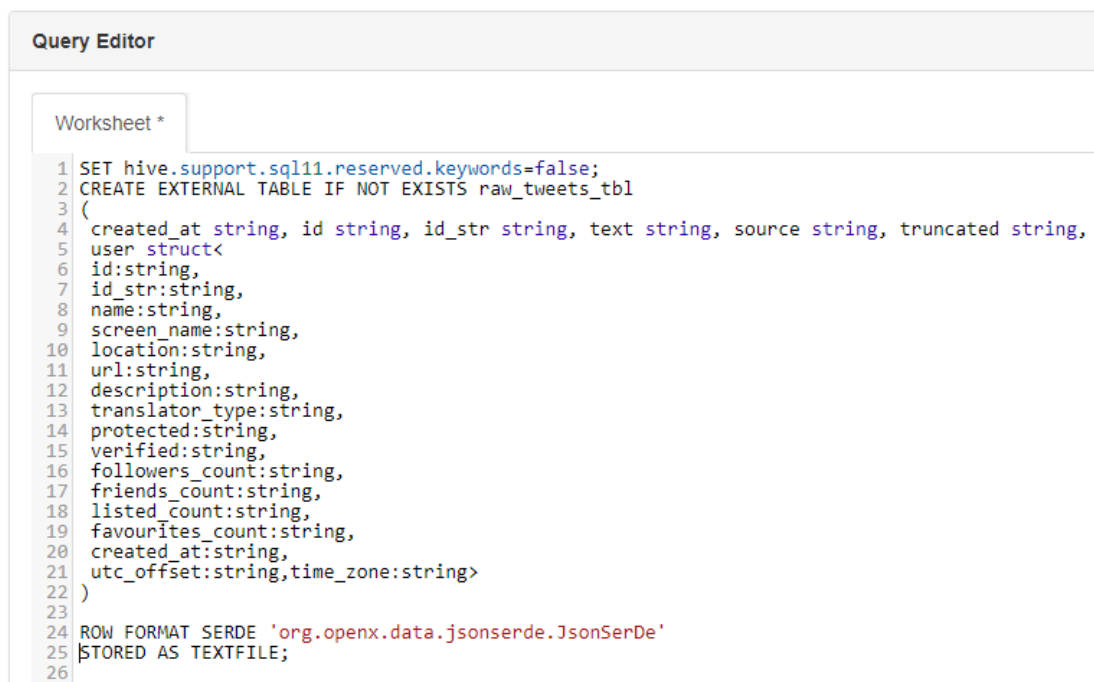
3.5. Create Hive table

Now we can create the hive table to load tweets in it. In the query below I have only selected some columns from the json file to show that we can just load the columns we require later in the process. Use the query or modify it to create the table where tweets will be loaded.

```
SET hive.support.sql11.reserved.keywords=false;
```

```
CREATE EXTERNAL TABLE IF NOT EXISTS <your-table-name>(created_at string,  
id string, id_str string, text string, source string, truncated string,  
user struct< id:string, id_str:string, name:string, screen_name:string,  
location:string, url:string,description:string, translator_type:string,  
protected:string, verified:string, followers_count:string,  
friends_count:string, listed_count:string, favourites_count:string,  
created_at:string, utc_offset:string,time_zone:string>)
```

```
ROW FORMAT SERDE 'org.openx.data.jsonserde.JsonSerDe' STORED AS TEXTFILE;
```



```
Query Editor  
Worksheet *  
1 SET hive.support.sql11.reserved.keywords=false;  
2 CREATE EXTERNAL TABLE IF NOT EXISTS raw_tweets_tbl  
3 (  
4   created_at string, id string, id_str string, text string, source string, truncated string,  
5   user struct<  
6     id:string,  
7     id_str:string,  
8     name:string,  
9     screen_name:string,  
10    location:string,  
11    url:string,  
12    description:string,  
13    translator_type:string,  
14    protected:string,  
15    verified:string,  
16    followers_count:string,  
17    friends_count:string,  
18    listed_count:string,  
19    favourites_count:string,  
20    created_at:string,  
21    utc_offset:string,time_zone:string>  
22 )  
23  
24 ROW FORMAT SERDE 'org.openx.data.jsonserde.JsonSerDe'  
25 STORED AS TEXTFILE;  
26
```

3.6. Load Data into Hive Table

Now we will load the tweets into the hive table we created above. Use the following command to load the tweets.

```
SET hive.support.sql11.reserved.keywords=false;
```

```
LOAD DATA INPATH '<path-to-tweets-folder/' OVERWRITE INTO TABLE <your-  
table-name>;
```

```
Query Editor

Worksheet *

1 SET hive.support.sql11.reserved.keywords=false;
2 LOAD DATA INPATH '/usr/tweet_received/'
3 OVERWRITE INTO TABLE raw_tweets_tbl;
```

3.7. Create processed table

The next step of the process is to load the table we just created into Pig and extract year, month, day, hour, minute, id and text. We will then store these into a new hive table which we need to create before going to Pig as we need a place to store the output of Pig. Use the following lines to create the table.

```
SET hive.support.sql11.reserved.keywords=false;
```

```
CREATE EXTERNAL TABLE IF NOT EXISTS <your-table-name>(year string, month  
string, day string, hour string, minute string, id string, text string)
```

```
ROW FORMAT SERDE 'org.openx.data.jsonserde.JsonSerDe' STORED AS TEXTFILE;
```

```
Query Editor

Worksheet *

1 SET hive.support.sql11.reserved.keywords=false;
2 CREATE EXTERNAL TABLE IF NOT EXISTS proc_tweets_tbl(
3   year string,
4   month string,
5   day string,
6   hour string,
7   minute string,
8   id string,
9   text string
10 )
11 ROW FORMAT SERDE 'org.openx.data.jsonserde.JsonSerDe'
12 STORED AS TEXTFILE;
13
```

3.8. Load Hive Table to Pig and Process it

Next we upload the hive table into pig and process it. We will extract only the data we need from the tweets and store it in the new table we just created. In addition, we will extract the created_at column and split it to get 5 separate columns of year, month, date, hour and minutes.

- Open the Pig View and create a new script. Give it the name you want and create it.

- Open the script and paste the following code in it changing the code where required.

```
REGISTER hdfs:///<serde-path>/<serde-name>.jar;

a = LOAD '<tweets-table-name>' USING
org.apache.hive.hcatalog.pig.HCatLoader();
f = FOREACH a GENERATE ToDate(created_at, 'EEE MMM dd HH:mm:ss Z yyyy') as
(date_time:DateTime ),id as iden,text as t;

y = FOREACH f GENERATE
GetYear(date_time)as (year:chararray),
GetMonth(date_time)as(month:chararray),
GetDay(date_time)as(day:chararray),
GetHour(date_time)as(hour:chararray),
GetMinute(date_time)as(minute:chararray),
iden as id,
t as text;

STORE y INTO '<processed-table-name>' USING
org.apache.hive.hcatalog.pig.HCatStorer();
```

The screenshot shows the 'kurome' web interface. At the top, there's a header with the 'kurome' logo and a pencil icon. Below the header, there are two tabs: 'PIG helper' and 'UDF helper'. On the right side, there's a file path indicator: '/tmp/.pigscrip...'. The main area displays a Pig script with line numbers from 1 to 14. The script is a Pig Latin script that registers a jar, loads data from a table, processes it with various date and time functions, and stores the result into another table. The script is as follows:

```
1 REGISTER hdfs:///usr/json-serde-1.1.9.9.jar;
2 a = LOAD 'raw_tweets_tbl' USING org.apache.hive.hcatalog.pig.HCatLoader();
3 f = FOREACH a GENERATE ToDate(created_at, 'EEE MMM dd HH:mm:ss Z yyyy') as
4 (date_time:DateTime ),id as iden,text as t;
5 y = FOREACH f GENERATE
6 GetYear(date_time)as (year:chararray),
7 GetMonth(date_time)as(month:chararray),
8 GetDay(date_time)as(day:chararray),
9 GetHour(date_time)as(hour:chararray),
10 GetMinute(date_time)as(minute:chararray),
11 iden as (id:chararray),
12 t as (text:chararray);
13 STORE y INTO 'proc_tweets_tbl' USING org.apache.hive.hcatalog.pig.HCatStorer();
14 |
```

- Scroll to the end of the page and add the argument `-useHCatalog` otherwise the HCatalog service will not be accessible.

-useHCatalog

+ Add

Arguments

-useHCatalog ×

Pig argument

+ Add

Version 2.0

- After this execute the script. You can check Execute on Tez to use the Tez execution engine. Otherwise it will run on MapReduce.

3.9. Remove new line characters in hive

The next step is to use Hive's explode function on the text column (explode function is explained in the next section). Before that, however, we need to remove all newline characters in the text column as they can cause unwanted behavior. Use the following query in Hive View to remove newline characters. **Make sure the Serde is added to the worksheet before executing this query.**

```
SET hive.support.sql11.reserved.keywords=false;
```

```
create view <view-name> as SELECT hour,id,regexp_replace(text,'\n','') as text
```

```
FROM <processed-table-name>;
```

Query Editor

Worksheet *

```
1 SET hive.support.sql11.reserved.keywords=false;
2 create view temp as SELECT hour,id,regexp_replace(text,'\n','') as text
3 FROM proc_tweets_tbl
```

3.10. Hive Explode Function

Next we will use hive's explode function. The explode function separate the words in a sentence into multiple rows. For example, if a row contains a sentence that has 5 words then then the explode

function will create 5 rows with each word in the sentence in a different row. The other columns in the table will be replicated as they are in the 5 rows.

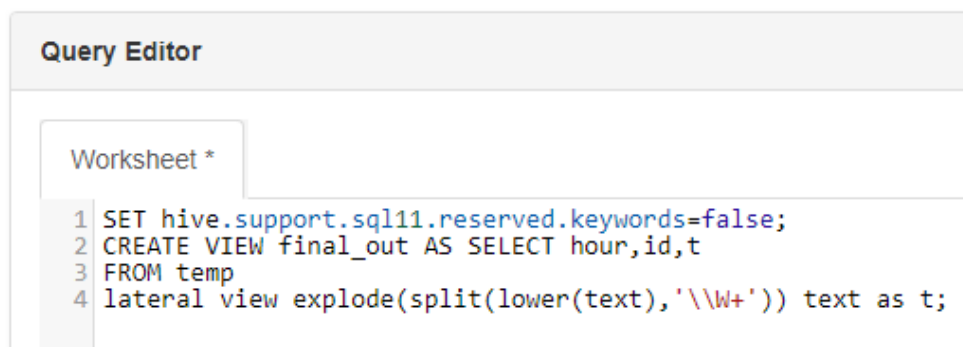
Use the following query to execute the explode function

```
SET hive.support.sql11.reserved.keywords=false;

CREATE VIEW <view-name> AS SELECT hour,id,t

FROM <view-name-created-in-last-step>

lateral view explode(split(lower(text),'\W+')) text as t;
```



The screenshot shows a 'Query Editor' window with a 'Worksheet *' tab. It contains a SQL query with four lines: setting a Hive configuration, creating a view, selecting from a temporary table, and using a lateral view with the explode function to split text into words.

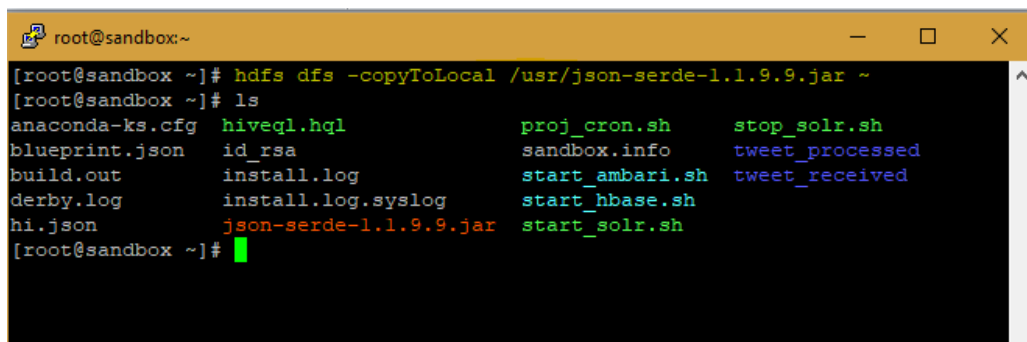
```
1 SET hive.support.sql11.reserved.keywords=false;
2 CREATE VIEW final_out AS SELECT hour,id,t
3 FROM temp
4 lateral view explode(split(lower(text),'\W+')) text as t;
```

3.11. Spark Group By Function

After separating the words we will now take the data into spark and get insights from the data. Right now the following query will get us the **number of words used per hour** in these tweets. We can write more complex queries in spark and even create graphs from this data using Spark's GraphX library but that will be for another time.

- Go to the shell and enter the following command to get the serde into the ndfs. This step is done because the Spark only adds jars from the ndfs so we need to put the jar file in the ndfs.

```
hdfs dfs -copyToLocal /usr/json-serde-1.1.9.9.jar ~
```

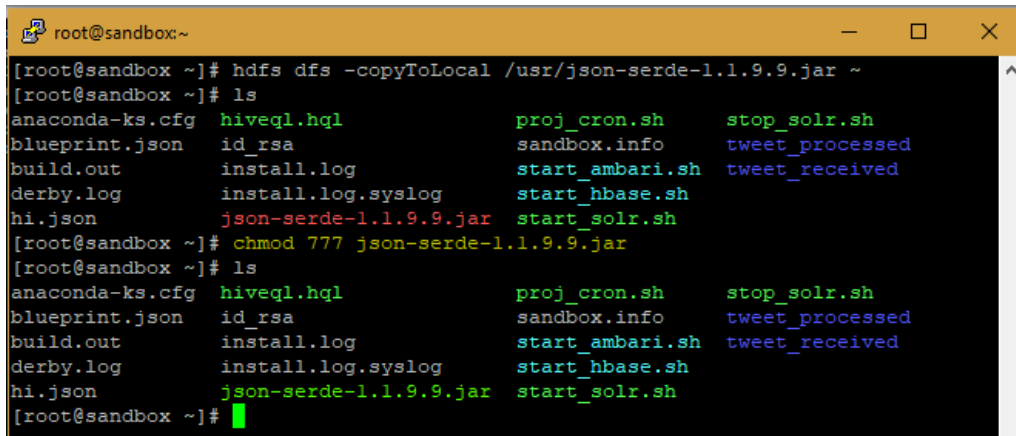


The screenshot shows a terminal window with the prompt 'root@sandbox:~'. The user enters the command 'hdfs dfs -copyToLocal /usr/json-serde-1.1.9.9.jar ~'. The output shows a list of files in the current directory, including 'anaconda-ks.cfg', 'blueprint.json', 'build.out', 'derby.log', 'hi.json', 'hiveql.hql', 'id_rsa', 'install.log', 'install.log.syslog', 'json-serde-1.1.9.9.jar', 'proj_cron.sh', 'sandbox.info', 'start_ambari.sh', 'start_hbase.sh', 'start_solr.sh', 'stop_solr.sh', and 'tweet_processed'.

```
root@sandbox:~# hdfs dfs -copyToLocal /usr/json-serde-1.1.9.9.jar ~
root@sandbox:~# ls
anaconda-ks.cfg  hiveql.hql          proj_cron.sh       stop_solr.sh
blueprint.json   id_rsa              sandbox.info       tweet_processed
build.out        install.log          start_ambari.sh    tweet_received
derby.log        install.log.syslog   start_hbase.sh
hi.json          json-serde-1.1.9.9.jar start_solr.sh
```

- Give execution permission to the serde

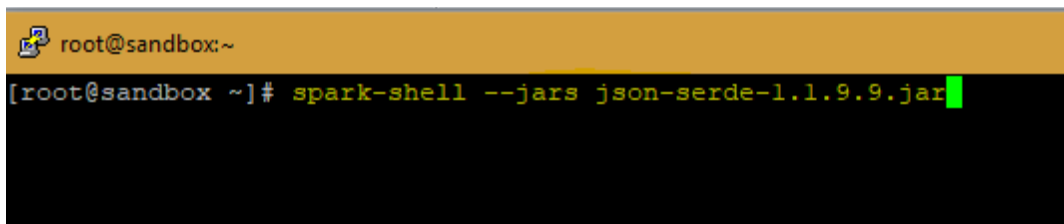
```
chmod 777 json-serde-1.1.9.9.jar
```



```
root@sandbox:~  
[root@sandbox ~]# hdfs dfs -copyToLocal /usr/json-serde-1.1.9.9.jar ~  
[root@sandbox ~]# ls  
anaconda-ks.cfg  hiveql.hql          proj_cron.sh       stop_solr.sh  
blueprint.json  id_rsa              sandbox.info       tweet_processed  
build.out        install.log          start_ambari.sh    tweet_received  
derby.log        install.log.syslog   start_hbase.sh  
hi.json          json-serde-1.1.9.9.jar start_solr.sh  
[root@sandbox ~]# chmod 777 json-serde-1.1.9.9.jar  
[root@sandbox ~]# ls  
anaconda-ks.cfg  hiveql.hql          proj_cron.sh       stop_solr.sh  
blueprint.json  id_rsa              sandbox.info       tweet_processed  
build.out        install.log          start_ambari.sh    tweet_received  
derby.log        install.log.syslog   start_hbase.sh  
hi.json          json-serde-1.1.9.9.jar start_solr.sh  
[root@sandbox ~]#
```

- Start the spark shell with the following command to add the serde to the spark classpath

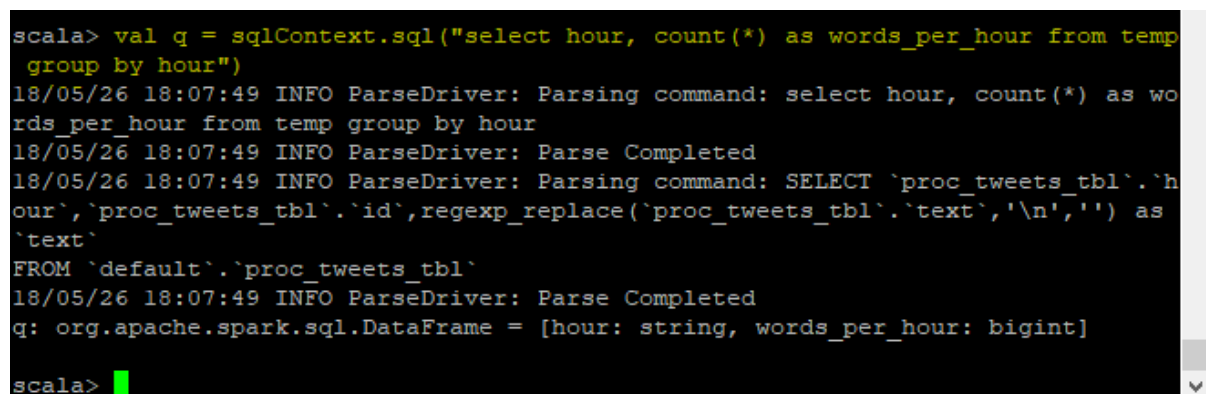
```
spark-shell --jars json-serde-1.1.9.9.jar
```



```
root@sandbox:~  
[root@sandbox ~]# spark-shell --jars json-serde-1.1.9.9.jar
```

- After the shell has started enter the following commands to load the hive view, process it and store the result in a new hive table.

```
val q = sqlContext.sql("select hour, count(*) as words_per_hour from  
final_out group by hour")
```

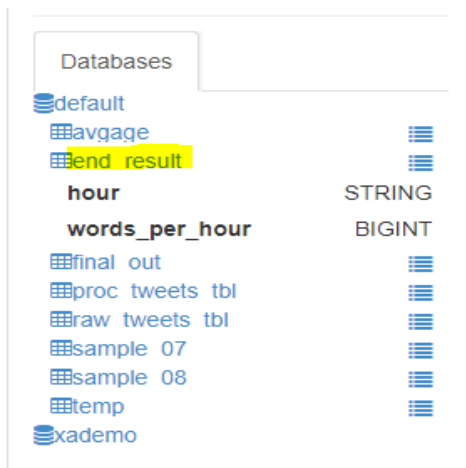


```
scala> val q = sqlContext.sql("select hour, count(*) as words_per_hour from temp  
group by hour")  
18/05/26 18:07:49 INFO ParseDriver: Parsing command: select hour, count(*) as wo  
rds_per_hour from temp group by hour  
18/05/26 18:07:49 INFO ParseDriver: Parse Completed  
18/05/26 18:07:49 INFO ParseDriver: Parsing command: SELECT `proc_tweets_tbl`.`h  
our`, `proc_tweets_tbl`.`id`, regexp_replace(`proc_tweets_tbl`.`text`, '\n', '') as  
`text`  
FROM `default`.`proc_tweets_tbl`  
18/05/26 18:07:49 INFO ParseDriver: Parse Completed  
q: org.apache.spark.sql.DataFrame = [hour: string, words_per_hour: bigint]  
scala>
```

```
q.saveAsTable("end_result")
```

```
scala> q.saveAsTable("end_result")
```

- After the table has been saved go to the Hive view through Ambari and see the created table.



Databases	
default	
avgage	
end_result	
hour	STRING
words_per_hour	BIGINT
final out	
proc tweets tbl	
raw tweets tbl	
sample 07	
sample 08	
temp	
xademo	

- Contents of the table:

end_result.hour	end_result.words_per_hour
5	12
6	10
7	15
9	11
10	20
11	34
12	45
13	42
14	40
15	30
16	6